

eZ Publish and Varnish

Web Summer Camp 2017

About me

- Hrvoje Knežević
 - Working at Netgen since 2012
 - Senior Symfony / eZ Publish developer, project tech lead
- <https://twitter.com/hknezevic>
- <https://www.linkedin.com/in/hknezevic/>
- <https://gitlab.com/hknezevic>
- <https://github.com/hknezevic>
- email: hrvoje.k@netgen.hr

...before we start...

- `$ cd /var/www/html/summercamp/ezhttpcaching`
- `$ git pull`
- `$./prepare.sh`

Varnish

- HTTP proxy
- HTTP accelerator for content-heavy dynamic web sites
- stores data in virtual memory, streams log information in shared memory
- heavily threaded

VCL

- *Varnish Configuration Language*
- domain-specific language, simple syntax similar to C (operators) or script languages (regex)
- translated to C, compiled as a shared object
- subroutines, ACL

Functions, actions and variables

- https://book.varnish-software.com/4.0/chapters/VCL_Basics.html
- <https://varnish-cache.org/docs/4.0/reference/vcl.html>
- <https://varnish-cache.org/docs/4.0/users-guide/vcl-actions.html>
- <https://www.varnish-cache.org/docs/trunk/users-guide/vcl-built-in-subs.html>

Functions, actions and variables

- vcl_recv, vcl_fetch
- pass, deliver, pipe, hash
- **req**, bereq, **resp**, beresp, obj(.ttl/.grace, ...)

Varnish configuration file

- <https://github.com/ezsystems/ezplatform/blob/master/doc/varnish/vcl/varnish4.vcl>

X-User-Hash

- User context (additional Vary header)
- In the context of eZ Platform, X-User-Hash is generated based on the roles defined for user and user groups
- Anonymous user has hard-coded hash in VCL - prevents additional backend requests to retrieve user hash

```
$ curl -I -H "Accept: application/vnd.fos.user-context-hash" http://  
ezhttpcaching.websc/_fos_user_context_hash
```

Tools

- **varnishlog**
 - displays formatted log stream from shared memory
 - <https://www.varnish-cache.org/docs/5.0/reference/varnishlog.html>
 - <https://www.varnish-cache.org/trac/wiki/VarnishlogExamples>

Tools

- **varnishadm**
 - CLI tool for accessing Varnish admin console

```
varnishadm "ban req.http.host ~ ezhttpcaching.websc"
```

```
varnishadm "ban req.http.host == ezhttpcaching.websc && req.url == /Projects"
```

```
varnishadm "ban req.http.host ~ ezhttpcaching.websc && req.url ~ .css$"
```

```
varnishadm "ban obj.http.X-ContentType-Identifier == blog_post"
```

Varnish and eZ

- <https://doc.ez.no/display/EZP/Using+Varnish>
- Varnish is used as an HTTP cache
- comparison to eZ Publish legacy
 - “replaces” legacy view cache
 - ESI fragments “replace” legacy block cache

Configuration

- apache/htaccess :

```
SetEnv SYMFONY_ENV prod
```

```
SetEnv SYMFONY_HTTP_CACHE 0
```

```
SetEnv SYMFONY_TRUSTED_PROXIES "127.0.0.1"
```

- nginx:

```
fastcgi_param SYMFONY_ENV prod;
```

```
fastcgi_param SYMFONY_HTTP_CACHE 0;
```

```
fastcgi_param SYMFONY_TRUSTED_PROXIES "127.0.0.1";
```

Configuration

- ezplatform.yml:

```
ezpublish:
  http_cache:
    purge_type: http
  system:
    site_group:
      http_cache:
        purge_servers: [http://127.0.0.1]
  default:
    content:
      view_cache: true
      ttl_cache: true
      default_ttl: 3600
```

Edge-Side Includes (ESI)

- small markup language for edge level dynamic web content assembly
- HTML statements over HTTP for including fragments of web pages in other web pages
- `<esi:include src=.../>`

Debugging ESI

```
$ php app/console cache:clear --no-debug
```

```
$ curl -H "Surrogate-Capability: abc=ESI/1.0" http://ezhttpcaching.websc
```


Edge-side includes

- eZ content views:

```
{{ render_ez( controller(
    "ez_content:viewAction", {"contentId": 123, "viewType": "line"}
)) }}
```

- Custom Symfony controllers:

```
{{ render_ez(controller("AppBundle:Parts:getLatestBlogPosts")) }}
```

Edge-side includes

- `Symfony\Component\HttpFoundation\Response`
 - `setPublic()`
 - `setPrivate()`
 - `setMaxAge()`
 - `setSharedMaxAge()`
 - `setVary()`
 - `headers`

Tagging view responses with custom headers

- eZ sets X-Location-Id header by default on every view_content response
- <https://github.com/eZsystems/ezpublish-kernel/blob/master/eZ/Bundle/EzPublishCoreBundle/EventListener/CacheViewResponseListener.php>
- we can implement our own response listeners for tagging content with additional headers
 - tag: kernel.event_subscriber

Smart view cache clearing

- provides possibility to clear cache for locations or content that is in relation with the content being currently cleared
 - tag: `ezpublish.http_cache.event_subscriber`
- <https://github.com/ezsystems/ezpublish-kernel/tree/master/eZ/Publish/Core/MVC/Symfony/Cache/Http/EventListener>
- extend list of location IDs which will be invalidated

Clearing cache by custom header

- We can utilise the fos_http_cache.cache_manager to manually invalidate HTTP cache
 - <https://github.com/FriendsOfSymfony/FOSHttpCacheBundle/blob/master/src/CacheManager.php>
 - <https://github.com/FriendsOfSymfony/FOSHttpCache/blob/master/src/CacheInvalidator.php>

Legacy tip

- If you are using the eZ Platform with legacy support, make use of the legacy smart view cache!
- Legacy view cache clearing triggers invalidation of the HTTP cache as well

Pro tip for the multi-site setups

- Varnish tags all of the responses with the X-Host header by default
- You can override the default FOSPurgeClient and apply additional headers on the BAN operations
 - <https://github.com/ezsystems/ezpublish-kernel/blob/master/eZ/Publish/Core/MVC/Symfony/Cache/Http/FOSPurgeClient.php>

What's next?

- platform-http-cache
 - <https://github.com/ezsystems/ezplatform-http-cache>
 - Experimental HTTP cache handling for eZ Platform.
 - Utilizes xkey module for Varnish